

Module V:

Graphs and Sorting

I. Graphs

Graphs are non-linear data structures used to represent connections or relationships between discrete objects. They consist of a set of **vertices** (also called nodes) and a set of **edges** (also called links or arcs) that connect pairs of vertices. Graphs are incredibly versatile and find applications in diverse fields, including social networks, transportation systems, computer networks, and more. To work with graphs programmatically, we need efficient ways to represent them in computer memory. The two most common representation methods are the Adjacency Matrix and the Adjacency List.

1. Adjacency Matrix:

An adjacency matrix is a square 2D array, typically denoted as A , where the dimensions are $V \times V$, with V being the number of vertices in the graph. Each cell $A[i][j]$ in the matrix stores information about the connection between vertex i and vertex j .

- For an **unweighted graph**: $A[i][j]$ is 1 if there is an edge between vertex i and vertex j , and 0 otherwise.
- For a **weighted graph**: $A[i][j]$ stores the weight of the edge between vertex i and vertex j , and infinity or a special value (like 0, depending on context) if there is no direct edge.
- For an **undirected graph**: The adjacency matrix is symmetric, meaning $A[i][j]=A[j][i]$.
- For a **directed graph**: $A[i][j]$ is 1 (or weight) if there is an edge from i to j , and 0 otherwise. $A[j][i]$ would represent an edge from j to i .

Advantages of Adjacency Matrix:

- **Easy to implement and understand.**
- **Quickly check for an edge between two specific vertices**: $O(1)$ time complexity by checking $A[i][j]$.
- **Adding/deleting an edge is efficient**: $O(1)$ time.

Disadvantages of Adjacency Matrix:

- **Space Inefficient for Sparse Graphs**: For graphs with many vertices but relatively few edges (sparse graphs), the matrix will contain mostly zeros, leading to wasted space ($O(V^2)$ space complexity).
- **Iterating over neighbors is inefficient**: To find all neighbors of a vertex V , you need to iterate through an entire row (or column) of size V , which takes $O(V)$ time.

2. Adjacency List:

An adjacency list is a collection of linked lists or dynamic arrays, one for each vertex in the graph. For each vertex V_i , its corresponding list contains all vertices V_j such that there is an edge from V_i to V_j .

- For an **unweighted graph**: Each list stores the neighboring vertices.
- For a **weighted graph**: Each list stores pairs of (neighboring vertex, edge weight).
- For an **undirected graph**: If there's an edge between i and j , then j will be in i 's list, and i will be in j 's list.
- For a **directed graph**: If there's an edge from i to j , then j will be in i 's list. i will not necessarily be in j 's list.

Advantages of Adjacency List:

- **Space Efficient for Sparse Graphs**: It only stores existing edges, making it much more space-efficient for sparse graphs ($O(V+E)$ space complexity, where E is the number of edges).
- **Efficient for finding neighbors**: To find all neighbors of a vertex, you just iterate through its corresponding list, which takes $O(\text{degree of vertex})$ time.

Disadvantages of Adjacency List:

- **Checking for an edge is less efficient**: To check if an edge exists between two specific vertices i and j , you might need to iterate through i 's adjacency list, which can take $O(\text{degree of } i)$ time in the worst case.
- **Slightly more complex to implement than an adjacency matrix**.

The choice between an adjacency matrix and an adjacency list depends on the specific graph characteristics and the operations that will be performed most frequently. For dense graphs (many edges), an adjacency matrix might be suitable. For sparse graphs (few edges), an adjacency list is generally preferred due to its space efficiency.

BFS and DFS

Breadth-First Search (BFS) and Depth-First Search (DFS) are two fundamental graph traversal algorithms. They provide systematic ways to explore all the vertices and edges in a graph, starting from a given source vertex. While both algorithms explore the graph, they do so in different orders, leading to different applications and characteristics.

1. Breadth-First Search (BFS):

BFS explores a graph level by level. It starts at a source node, visits all its immediate neighbors, then visits all the neighbors of those neighbors, and so on, moving outward in layers. BFS uses a queue data structure to manage the vertices to be visited.

Algorithm:

1. Start with a source vertex, mark it as visited, and add it to a queue.
2. While the queue is not empty:
 - a. Dequeue a vertex v .
 - b. Process v (e.g., print it).
 - c. For each unvisited neighbor u of v :
 - i. Mark u as visited.
 - ii. Enqueue u .

Characteristics and Applications of BFS:

- **Shortest Path in Unweighted Graphs:** BFS guarantees finding the shortest path (in terms of number of edges) from the source to all reachable vertices in an unweighted graph. This is because it explores vertices in increasing order of their distance from the source.
- **Connectivity:** Can be used to determine if a graph is connected or to find connected components.
- **Minimum Spanning Tree (MST) for Unweighted Graphs:** While not directly an MST algorithm, its level-by-level exploration can be used as a building block for some MST approaches.
- **Finding Cycles:** Can detect cycles in undirected graphs.
- **Web Crawlers:** Used to traverse web pages linked from a starting page.

2. Depth-First Search (DFS):

DFS explores a graph by going as deep as possible along each branch before backtracking. It starts at a source node, explores one of its unvisited neighbors, then explores that neighbor's unvisited neighbors, and so on, until it reaches a dead end (a node with no unvisited neighbors). At this point, it backtracks to the most recent node that has unvisited neighbors and continues the exploration. DFS typically uses a stack data structure (implicitly via recursion or explicitly using an iterative approach).

Algorithm (Recursive):

1. Start with a source vertex, mark it as visited.
2. Process the source vertex.
3. For each unvisited neighbor u of the source vertex:
a. Recursively call DFS on u .

Characteristics and Applications of DFS:

- **Cycle Detection:** Efficiently detects cycles in both directed and undirected graphs.
- **Topological Sorting:** Used for ordering tasks with dependencies in directed acyclic graphs (DAGs).
- **Strongly Connected Components (SCCs):** Algorithms like Tarjan's or Kosaraju's for finding SCCs in directed graphs are based on DFS.
- **Path Finding:** Can find *any* path between two nodes (not necessarily the shortest).
- **Solving Mazes:** A natural fit for exploring paths in a maze.
- **Game AI:** Used in game tree traversal (e.g., to find winning moves).

Comparison:

Feature	BFS	DFS
Data Structure	Queue	Stack (or Recursion)

Traversal Order	Level by Level (Breadth-wise)	As deep as possible (Depth-wise)
Shortest Path	Guaranteed for unweighted graphs	Not guaranteed
Memory Usage	Can be high for wide graphs	Can be high for deep graphs
Recursion	Typically iterative	Often recursive

Both BFS and DFS have a time complexity of $O(V+E)$ for graphs represented using an adjacency list, or $O(V^2)$ for graphs represented using an adjacency matrix, where V is the number of vertices and E is the number of edges. The choice between BFS and DFS depends on the specific problem being solved, as their different exploration strategies are better suited for different tasks.

II. Hashing Techniques

Hashing is a fundamental computer science concept used to efficiently store and retrieve data. It involves mapping data of arbitrary size (the "key") to a fixed-size value (the "hash value" or "hash code"). This hash value then serves as an index into an array or table, called a **hash table** (or hash map), where the actual data (or a reference to it) is stored. The primary goal of hashing is to achieve average $O(1)$ time complexity for insertion, deletion, and lookup operations, significantly faster than the $O(\log n)$ for balanced trees or $O(n)$ for unsorted arrays.

The two main components of a hashing system are:

1. **Hash Function:** This is an algorithm that takes a key as input and returns a hash value.

A good hash function should:

- Be deterministic: The same key always produces the same hash value.
- Be efficient to compute.
- Distribute keys uniformly across the hash table to minimize collisions.
- Minimize collisions: Different keys should ideally produce different hash values.

2. **Hash Table (or Hash Map):** This is the data structure that stores the key-value pairs.

It's essentially an array where each index corresponds to a hash value.

Collisions:

A collision occurs when two different keys produce the same hash value. Since the hash table has a fixed size, collisions are unavoidable as the number of possible keys is often much larger than the number of available slots in the table. Effective collision resolution strategies are crucial for the performance of a hash table.

Collision Resolution Techniques:

1. Separate Chaining:

- Each slot (or "bucket") in the hash table is a pointer to a linked list (or another data structure like a small dynamic array or even another hash table for very large datasets).
- When a collision occurs, the new key-value pair is simply added to the linked list at the calculated hash index.
- **Advantages:** Simple to implement, handles a large number of collisions gracefully, and the hash table never "fills up."
- **Disadvantages:** Requires extra space for pointers, and performance degrades if linked lists become very long (effectively $O(N)$ in worst case if all keys hash to same bucket).

2. Open Addressing (Probing):

- Instead of using separate data structures, all key-value pairs are stored directly within the hash table array.

- When a collision occurs, the algorithm "probes" for an alternative empty slot in the table using a predefined sequence.
- Common probing methods:
 - **Linear Probing:** If the primary hash index h is occupied, try $h+1, h+2, h+3$, and so on (modulo table size).
 - **Issue:** Can lead to **primary clustering**, where long runs of occupied slots form, increasing search times.
 - **Quadratic Probing:** If h is occupied, try $h+1^2, h+2^2, h+3^2$, etc. (modulo table size).
 - **Issue:** Can lead to **secondary clustering**, where keys that hash to the same initial location follow the same probe sequence.
 - **Double Hashing:** Uses a second hash function to determine the step size for probing. If the primary hash index h is occupied, the next probe location is $h + \text{step_size}$, $h + 2*\text{step_size}$, etc., where step_size is calculated by a second hash function.
 - **Advantages:** Reduces clustering significantly.
 - **Disadvantages:** Requires a second hash function.
- **Advantages of Open Addressing:** No extra space for pointers, better cache performance due to data locality.
- **Disadvantages of Open Addressing:** Table can fill up (requires resizing), deletion is complex (requires "tombstones" to mark deleted slots so searches don't stop prematurely), sensitive to load factor (ratio of occupied slots to total slots).

Load Factor:

The load factor (α) is a critical parameter for hash table performance, calculated as table size / number of elements.

- For **separate chaining**, a load factor greater than 1 is acceptable, though performance degrades as it increases.
- For **open addressing**, the load factor must be ≤ 1 . If it approaches 1, performance deteriorates rapidly, necessitating **resizing** (creating a larger table and re-hashing all existing elements into

III. Sorting Algorithms

Sorting is the process of arranging elements of a list or an array in a specific order, typically numerical or lexicographical (alphabetical). The goal is to organize data to make it easier to search, analyze, and process. Sorting algorithms are fundamental to computer science and have a wide range of applications, from database management and search engines to computational biology and data visualization. While many sorting algorithms exist, we will focus on five commonly taught and distinct methods: Bubble Sort, Selection Sort, Insertion Sort, Merge Sort, and Quick Sort. These algorithms vary significantly in their approach, efficiency, and practical applicability.

The efficiency of a sorting algorithm is typically measured by its time complexity (how the execution time grows with the input size n) and space complexity (how much auxiliary memory it requires). Sorting algorithms can be classified based on various factors:

- **Time Complexity:** How many comparisons and swaps are needed.
- **Space Complexity:** Whether they sort "in-place" (requiring minimal additional memory, $O(1)$ or $O(\log n)$) or require significant auxiliary space ($O(n)$).
- **Stability:** A sorting algorithm is stable if it preserves the relative order of equal elements. For example, if two elements A and B have the same value, and A appears before B in the original list, a stable sort will ensure A still appears before B in the sorted list.
- **Adaptivity:** An adaptive algorithm takes advantage of existing order in the input. If the input is already partially sorted, an adaptive algorithm might perform better.

Let's delve into each of the specified algorithms.

1. Bubble Sort

Bubble Sort is one of the simplest sorting algorithms, but also one of the least efficient for large datasets. It repeatedly steps through the list, compares adjacent elements, and swaps them if they are in the wrong order. The pass through the list is repeated until no swaps are needed, which indicates that the list is sorted. The name "bubble sort" comes from the way smaller (or larger, depending on the sort order) elements "bubble" to their correct positions at the beginning (or end) of the list.

How it Works (Ascending Order):

1. **Comparison and Swap:** Start from the beginning of the list. Compare the first two elements. If the first is greater than the second, swap them.
2. **Move to Next Pair:** Move to the next pair of elements (second and third), compare, and swap if necessary.
3. **Iterate Through Pass:** Continue this process until the end of the list is reached. After one complete pass, the largest element will have "bubbled" to its correct position at the very end of the unsorted portion of the list.
4. **Repeat Passes:** Repeat the entire process for the remaining unsorted portion of the list. In each subsequent pass, the range of elements to consider decreases by one, as the largest element of the current pass is correctly positioned.

5. **Termination:** The algorithm terminates when a full pass is completed without performing any swaps, indicating that the array is sorted.

Example (Ascending Order):

Consider the array: [5, 1, 4, 2, 8]

Pass 1:

- (5, 1) -> (1, 5) [1, 5, 4, 2, 8] (Swapped)
- (5, 4) -> (1, 4, 5, 2, 8) (Swapped)
- (5, 2) -> (1, 4, 2, 5, 8) (Swapped)
- (5, 8) -> (1, 4, 2, 5, 8) (No swap) After Pass 1: [1, 4, 2, 5, 8]. The largest element, 8, is now at its correct position.

Pass 2: (Considering [1, 4, 2, 5])

- (1, 4) -> (1, 4, 2, 5, 8) (No swap)
- (4, 2) -> (1, 2, 4, 5, 8) (Swapped)
- (4, 5) -> (1, 2, 4, 5, 8) (No swap) After Pass 2: [1, 2, 4, 5, 8]. The second largest, 5, is now at its correct position.

Pass 3: (Considering [1, 2, 4])

- (1, 2) -> (1, 2, 4, 5, 8) (No swap)
- (2, 4) -> (1, 2, 4, 5, 8) (No swap) After Pass 3: [1, 2, 4, 5, 8]. The third largest, 4, is now at its correct position.

Pass 4: (Considering [1, 2])

- (1, 2) -> (1, 2, 4, 5, 8) (No swap) After Pass 4: [1, 2, 4, 5, 8]. No swaps occurred in this pass. The list is sorted.

Complexity Analysis:

- **Time Complexity:**
 - **Worst Case:** $O(n^2)$ - Occurs when the array is in reverse order. In each pass, comparisons and swaps are performed roughly n times. There are $n-1$ passes. So, $(n-1)+(n-2)+\dots+1=O(n^2)$ comparisons and swaps.
 - **Average Case:** $O(n^2)$ - Similar to the worst case.
 - **Best Case:** $O(n)$ - Occurs when the array is already sorted. The algorithm will complete one pass, find no swaps, and terminate. This is an adaptive behavior.
- **Space Complexity:** $O(1)$ - It sorts in-place, requiring only a constant amount of extra space for temporary variables during swaps.
- **Stability:** Stable - It only swaps adjacent elements, preserving the relative order of equal elements.

Advantages:

- Simple to understand and implement.
- Efficient for very small lists or nearly sorted lists (best case).

Disadvantages:

- Very inefficient for large lists due to its quadratic time complexity.
- Many swaps can make it slow (especially compared to algorithms like Selection Sort which minimizes swaps).

Use Cases:

Bubble Sort is rarely used in practical applications for large datasets due to its poor performance. It is primarily used as a pedagogical tool to introduce the concept of sorting algorithms due to its simplicity.

2. Selection Sort

Selection Sort is another simple, intuitive, but generally inefficient sorting algorithm. It works by repeatedly finding the minimum (or maximum) element from the unsorted part of the list and putting it at the beginning (or end) of the sorted part.

How it Works (Ascending Order):

1. **Find Minimum:** In the first pass, scan the entire unsorted list to find the smallest element.
2. **Swap with First:** Swap this smallest element with the element at the first position of the unsorted list. Now, the first element is sorted.
3. **Repeat for Remaining:** Consider the remaining unsorted portion of the list (from the second element onwards). Find the smallest element in this new unsorted portion.
4. **Swap with Second:** Swap this found smallest element with the element at the second position of the original list.
5. **Continue:** Repeat this process for $n-1$ passes. After $n-1$ passes, the first $n-1$ elements will be in their correct sorted positions, and the last element will automatically be in its correct place.

Example (Ascending Order):

Consider the array: [64, 25, 12, 22, 11]

Pass 1 (Find minimum in [64, 25, 12, 22, 11]):

- Smallest element is 11.
- Swap 11 with 64 (at index 0).
- Array becomes: [11, 25, 12, 22, 64]
- Sorted part: [11]

Pass 2 (Find minimum in [25, 12, 22, 64]):

- Smallest element is 12.
- Swap 12 with 25 (at index 1).
- Array becomes: [11, 12, 25, 22, 64]
- Sorted part: [11, 12]

Pass 3 (Find minimum in [25, 22, 64]):

- Smallest element is 22.
- Swap 22 with 25 (at index 2).
- Array becomes: [11, 12, 22, 25, 64]
- Sorted part: [11, 12, 22]

Pass 4 (Find minimum in [25, 64]):

- Smallest element is 25.
- Swap 25 with 25 (at index 3). (No actual change)
- Array becomes: [11, 12, 22, 25, 64]
- Sorted part: [11, 12, 22, 25]

The array is now sorted: [11, 12, 22, 25, 64].

Complexity Analysis:

- **Time Complexity:**
 - **Worst Case:** $O(n^2)$ - In the first pass, $n-1$ comparisons are made. In the second, $n-2$, and so on. The total number of comparisons is $(n-1)+(n-2)+\dots+1=O(n^2)$. The number of swaps is always $n-1$.
 - **Average Case:** $O(n^2)$ - Consistent across various inputs.
 - **Best Case:** $O(n^2)$ - Even if the array is already sorted, it still performs all comparisons to find the minimum in each pass. It is not adaptive.
- **Space Complexity:** $O(1)$ - It sorts in-place, requiring only constant extra space for temporary variables during swaps.
- **Stability:** Unstable - It can change the relative order of equal elements because it swaps elements over long distances. For example, if you have [5a, 3, 5b] and 5b is swapped with 3 to place 3 in position, 5a and 5b might change relative order if 5b was the minimum in a later sub-array.

Advantages:

- Simple to understand and implement.
- Performs a minimal number of swaps ($n-1$ swaps in total), which can be beneficial in scenarios where writing to memory is an expensive operation.

Disadvantages:

- Inefficient for large lists due to its quadratic time complexity.
- Not adaptive; performs the same number of comparisons regardless of the initial order.

Use Cases:

Similar to Bubble Sort, Selection Sort is rarely used for large real-world applications. Its main advantage (minimal swaps) might make it marginally better than Bubble Sort in very specific edge cases where memory writes are extremely costly, but generally, other algorithms are preferred.

3. Insertion Sort

Insertion Sort builds the final sorted array (or list) one item at a time. It is significantly more efficient than Bubble Sort and Selection Sort for small lists and partially sorted lists. The algorithm iterates through the input list, taking one element at a time and inserting it into its correct position within the already sorted part of the list.

How it Works (Ascending Order):

1. **Initial Sorted Part:** Assume the first element of the array is already sorted.
2. **Take Next Element:** Start with the second element. This is the "key" element to be inserted.
3. **Compare and Shift:** Compare the key element with the elements in the sorted part, moving from right to left. If an element in the sorted part is greater than the key, shift it one position to the right to make space for the key.
4. **Insert Key:** Continue shifting until an element smaller than or equal to the key is found, or the beginning of the sorted part is reached. Insert the key into the empty spot.
5. **Repeat:** Repeat steps 2-4 for the remaining unsorted elements until all elements have been inserted into their correct positions.

Example (Ascending Order):

Consider the array: [12, 11, 13, 5, 6]

- **Initial:** [12 | 11, 13, 5, 6] (12 is considered sorted)

Pass 1 (Insert 11):

- Key = 11. Compare 11 with 12. Since $11 < 12$, shift 12 to the right.
- Array becomes: [| 12, 13, 5, 6] (space for 11)
- Insert 11: [11, 12 | 13, 5, 6]
- Sorted part: [11, 12]

Pass 2 (Insert 13):

- Key = 13. Compare 13 with 12. Since $13 > 12$, no shift needed.
- Insert 13: [11, 12, 13 | 5, 6]
- Sorted part: [11, 12, 13]

Pass 3 (Insert 5):

- Key = 5. Compare 5 with 13. 5 < 13, shift 13.
- Array: [11, 12, | 13, 5, 6]
- Compare 5 with 12. 5 < 12, shift 12.
- Array: [11, | 12, 13, 5, 6]
- Compare 5 with 11. 5 < 11, shift 11.
- Array: [| 11, 12, 13, 5, 6]
- Insert 5: [5, 11, 12, 13 | 6]
- Sorted part: [5, 11, 12, 13]

Pass 4 (Insert 6):

- Key = 6. Compare 6 with 13. 6 < 13, shift 13.
- Array: [5, 11, 12, | 13, 6]
- Compare 6 with 12. 6 < 12, shift 12.
- Array: [5, 11, | 12, 13, 6]
- Compare 6 with 11. 6 < 11, shift 11.
- Array: [5, | 11, 12, 13, 6]
- Compare 6 with 5. 6 > 5, stop.
- Insert 6: [5, 6, 11, 12, 13]

The array is now sorted: [5, 6, 11, 12, 13].

Complexity Analysis:

- **Time Complexity:**
 - **Worst Case:** $O(n^2)$ - Occurs when the array is in reverse order. Each element needs to be compared with and potentially shifted past all elements in the sorted portion.
 - **Average Case:** $O(n^2)$ - Similar to worst case.
 - **Best Case:** $O(n)$ - Occurs when the array is already sorted. Each element is compared only once with the last element of the sorted part, and no shifts occur. This is an adaptive behavior.
- **Space Complexity:** $O(1)$ - It sorts in-place, requiring only constant extra space for the key element and loop variables.
- **Stability:** Stable - It maintains the relative order of equal elements because new elements are inserted *after* existing equal elements.

Advantages:

- Simple to implement.
- Efficient for small datasets.
- Very efficient for data that is already substantially sorted (adaptive).
- Stable.
- Performs well on data that is received in a stream, as it can sort the data as it arrives.

Disadvantages:

- Inefficient for large, unsorted lists due to its quadratic time complexity.

Use Cases:

- Sorting very small arrays (e.g., in other sorting algorithms as a base case for recursion, like in hybrid sorts).
- Sorting lists that are mostly sorted or have few elements out of place.
- Online sorting scenarios where elements arrive one by one and need to be maintained in a sorted order.

4. Merge Sort

Merge Sort is a highly efficient, comparison-based sorting algorithm based on the **divide and conquer** paradigm. It is a stable sort and has a guaranteed $O(n \log n)$ time complexity, making it much faster than the $O(n^2)$ algorithms for large datasets.

How it Works:

1. **Divide:** Recursively divide the unsorted list into n sublists, each containing just one element. A list with one element is considered sorted.
2. **Conquer (Merge):** Repeatedly merge sublists to produce new sorted sublists until there is only one sorted list remaining. The core operation is the "merge" step, where two sorted sublists are combined into one larger sorted list.

The Merge Step (Key Operation):

Given two sorted sublists, say left and right, the merge function combines them into a single sorted list.

- It maintains two pointers, one for the left list and one for the right list, both initially pointing to the beginning of their respective lists.
- It compares the elements pointed to by the two pointers. The smaller element is copied to a temporary merged array, and its corresponding pointer is advanced.
- This process continues until all elements from one of the sublists have been copied.
- Finally, any remaining elements from the other sublist are copied directly to the merged array.

Example (Ascending Order):

Consider the array: [38, 27, 43, 3, 9, 82, 10]

1. Divide (Recursive Breakdown):

[38, 27, 43, 3, 9, 82, 10]

/

[38, 27, 43, 3] [9, 82, 10]

/ \ /

[38, 27] [43, 3] [9, 82] [10]

/ \ / \ /

[38] [27] [43] [3] [9] [82] [10] (Base cases: single elements are sorted)

2. Conquer (Merge Back Up):

- Merge [38] and [27] -> [27, 38]
- Merge [43] and [3] -> [3, 43]
- Merge [9] and [82] -> [9, 82]
- [10] remains as is.
- Merge [27, 38] and [3, 43] -> [3, 27, 38, 43]
- Merge [9, 82] and [10] -> [9, 10, 82]
- Finally, Merge [3, 27, 38, 43] and [9, 10, 82] -> [3, 9, 10, 27, 38, 43, 82]

The array is now sorted.

Complexity Analysis:

• Time Complexity:

- **Worst Case:** $O(n\log n)$ - At each level of recursion, we perform $O(n)$ work (the merge step). The number of levels of recursion is $\log n$ (since we divide the list in half each time). So, $O(n\log n)$.
- **Average Case:** $O(n\log n)$ - Consistent across various inputs.
- **Best Case:** $O(n\log n)$ - Even if the array is already sorted, it still performs the same divide and merge operations. It is not adaptive.

- **Space Complexity:** $O(n)$ - Merge sort requires an auxiliary array of size n to perform the merge operation. This is its main drawback.
- **Stability:** Stable - The merge operation can be implemented to preserve the relative order of equal elements. When comparing elements, if two are equal, prioritize the one from the left sublist.

Advantages:

- Guaranteed $O(n\log n)$ time complexity in all cases (worst, average, best).
- Stable sort.

- Well-suited for external sorting (sorting data that doesn't fit into memory) because it accesses data sequentially.

Disadvantages:

- Requires $O(n)$ auxiliary space, which can be a concern for very large datasets or memory-constrained environments.
- More complex to implement than $O(n^2)$ algorithms.

Use Cases:

- Sorting large datasets where consistent performance is critical.
- External sorting.
- Parallel computing (easily parallelizable due to its divide-and-conquer nature).
- As a subroutine in other algorithms, like external merge sort.

5. Quick Sort

Quick Sort is another highly efficient, comparison-based sorting algorithm that also follows the **divide and conquer** paradigm. It is one of the most widely used sorting algorithms due to its excellent average-case performance. It is generally faster than Merge Sort in practice due to better cache performance and fewer data movements, despite having the same average-case time complexity.

How it Works:

1. **Choose Pivot:** Select an element from the array, called the "pivot." The choice of pivot heavily influences performance. Common pivot choices include:
 - The first element.
 - The last element.
 - A random element.
 - The median of three (first, middle, last).
2. **Partition:** Rearrange the elements in the array such that all elements less than the pivot come before it, and all elements greater than the pivot come after it. Elements equal to the pivot can go on either side (or be handled separately). After partitioning, the pivot element is in its final sorted position.
3. **Conquer (Recursion):** Recursively apply Quick Sort to the subarray of elements with values less than the pivot and the subarray of elements with values greater than the pivot.
4. **Combine:** No explicit combine step is needed as the partitioning places elements in their correct relative positions. The array becomes sorted when all subarrays are sorted.

The Partition Step (Key Operation):

A common partitioning scheme (Lomuto partition scheme) selects the last element as the pivot.

- Initialize an index i (smaller element index) to low - 1.
- Iterate through the subarray from low to high - 1 (excluding the pivot).
- If an element arr[j] is less than or equal to the pivot, increment i and swap arr[i] with arr[j].
- Finally, swap the pivot (arr[high]) with arr[i+1]. Return i+1 as the pivot's final position.

Example (Ascending Order):

Consider the array: [10, 80, 30, 90, 40, 50, 70]

Let's choose the last element (70) as the pivot.

1. Initial Call: QuickSort([10, 80, 30, 90, 40, 50, 70], 0, 6)

2. Partition (Pivot = 70):

- [10, 80, 30, 90, 40, 50, 70]
- Iterate j from 0 to 5. i starts at -1.
 - j=0, arr[0]=10 <= 70: i=0, swap arr[0] with arr[0]. Array: [10, 80, 30, 90, 40, 50, 70]
 - j=1, arr[1]=80 > 70: No swap.
 - j=2, arr[2]=30 <= 70: i=1, swap arr[1] (80) with arr[2] (30). Array: [10, 30, 80, 90, 40, 50, 70]
 - j=3, arr[3]=90 > 70: No swap.
 - j=4, arr[4]=40 <= 70: i=2, swap arr[2] (80) with arr[4] (40). Array: [10, 30, 40, 90, 80, 50, 70]
 - j=5, arr[5]=50 <= 70: i=3, swap arr[3] (90) with arr[5] (50). Array: [10, 30, 40, 50, 80, 90, 70]
- Swap pivot (70) with arr[i+1] (which is arr[4], 80).
- Array after partition: [10, 30, 40, 50, 70, 90, 80]
- Pivot (70) is now at index 4. partition_index = 4.

3. Recursive Calls:

- QuickSort([10, 30, 40, 50], 0, 3) (left subarray)
- QuickSort([90, 80], 5, 6) (right subarray)

... This process continues recursively until all subarrays have 0 or 1 element, at which point they are considered sorted. The final result is [10, 30, 40, 50, 70, 80, 90].

Complexity Analysis:

- **Time Complexity:**

- **Worst Case:** $O(n^2)$ - Occurs when the pivot selection consistently results in highly unbalanced partitions (e.g., always choosing the smallest or largest element as the pivot in an already sorted or reverse-sorted array). This leads to one subproblem of size $n-1$ and another of size 0.
- **Average Case:** $O(n \log n)$ - When pivots are chosen well (e.g., random pivot or median-of-three), the partitions are relatively balanced, leading to logarithmic recursion depth.
- **Best Case:** $O(n \log n)$ - Occurs when the pivot consistently divides the array into two roughly equal halves.
- **Space Complexity:**
 - **Worst Case:** $O(n)$ - In the worst case (unbalanced partitions), the recursion depth can be n , leading to $O(n)$ space for the call stack.
 - **Average Case:** $O(\log n)$ - For balanced partitions, the recursion depth is $\log n$.
- **Stability:** Unstable - The partitioning process can change the relative order of equal elements (e.g., when swapping elements past each other).

Advantages:

- Highly efficient in practice for large datasets due to its cache-friendly memory access patterns.
- Average-case time complexity is $O(n \log n)$, which is optimal for comparison sorts.
- In-place sorting (or nearly in-place, depending on partition scheme and recursive call stack).

Disadvantages:

- Worst-case $O(n^2)$ performance, which can be a problem in real-time systems if pivot selection is poor.
- Not stable.
- More complex to implement correctly than some other algorithms.

Use Cases:

- General-purpose sorting in many programming languages' standard libraries (e.g., C++ `std::sort` often uses Introsort, a hybrid of QuickSort, HeapSort, and InsertionSort).
- Situations where average-case performance is acceptable, and memory usage is a concern.
- Sorting large arrays in main memory.

Summary of Sorting Algorithms:

Algorithm	Time	Time	Time	Space	Stable	Adaptive
-----------	------	------	------	-------	--------	----------

	Complexity (Worst)	Complexity (Average)	Complexity (Best)	Complexity		
Bubble Sort	$O(n^2)$	$O(n^2)$	$O(n)$	$O(1)$	Yes	Yes
Selection Sort	$O(n^2)$	$O(n^2)$	$O(n^2)$	$O(1)$	No	No
Insertion Sort	$O(n^2)$	$O(n^2)$	$O(n)$	$O(1)$	Yes	Yes
Merge Sort	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$	$O(n)$	Yes	No
Quick Sort	$O(n^2)$	$O(n \log n)$	$O(n \log n)$	$O(\log n)$ (avg), $O(n)$ (worst)	No	No

Choosing the Right Algorithm:

- For **small arrays or nearly sorted arrays**, **Insertion Sort** is often a good choice due to its simplicity and $O(n)$ best-case performance.
- For **large datasets** where **guaranteed performance** ($O(n \log n)$) and **stability** are crucial, **Merge Sort** is preferred, despite its $O(n)$ space complexity.
- For **large datasets** where **average-case performance** is paramount and **memory is a concern** (especially for recursive stack space), **Quick Sort** is usually the fastest choice, though its worst-case $O(n^2)$ needs to be managed (e.g., by using a good pivot selection strategy or hybrid algorithms).
- **Bubble Sort** and **Selection Sort** are generally not used for practical applications due to their $O(n^2)$ performance in most cases, except for educational purposes or extremely small, specific scenarios.

In real-world applications, highly optimized **hybrid sorting algorithms** (like Introsort, used in many standard library `sort()` implementations) combine the strengths of different algorithms (e.g., Quicksort for general cases, Heapsort for worst-case fallback, and Insertion Sort for small subarrays) to provide excellent performance across various input types.